

C3_W2_Assignment

January 25, 2021

1 Assignment 2: Deep N-grams

Welcome to the second assignment of course 3. In this assignment you will explore Recurrent Neural Networks RNN. - You will be using the fundamentals of google's [trax](#) package to implement any kind of deeplearning model.

By completing this assignment, you will learn how to implement models from scratch: - How to convert a line of text into a tensor - Create an iterator to feed data to the model - Define a GRU model using trax - Train the model using trax - Compute the accuracy of your model using the perplexity - Predict using your own model

1.1 Outline

- Section ??
- Section ??
 - Section ??
 - Section ??
 - * Section ??
 - Section ??
 - * Section ??
 - Section ??
- Section ??
 - Section ??
- Section ??
 - Section ??
 - * Section ??
- Section ??
 - Section ??
 - * Section ??
- Section ??

- Section ??

Overview

Your task will be to predict the next set of characters using the previous characters. - Although this task sounds simple, it is pretty useful. - You will start by converting a line of text into a tensor - Then you will create a generator to feed data into the model - You will train a neural network in order to predict the new set of characters of defined length. - You will use embeddings for each character and feed them as inputs to your model. - Many natural language tasks rely on using embeddings for predictions. - Your model will convert each character to its embedding, run the embeddings through a Gated Recurrent Unit GRU, and run it through a linear layer to predict the next set of characters.

The figure above gives you a summary of what you are about to implement. - You will get the embeddings; - Stack the embeddings on top of each other; - Run them through two layers with a relu activation in the middle; - Finally, you will compute the softmax.

To predict the next character: - Use the softmax output and identify the word with the highest probability. - The word with the highest probability is the prediction for the next word.

```
In [ ]: import os
import trax
import trax.fastmath.numpy as np
import pickle
import numpy
import random as rnd
from trax import fastmath
from trax import layers as tl

# set random seed
trax.supervised.trainer_lib.init_random_number_generators(32)
rnd.seed(32)
```

Part 1: Importing the Data

1.1 Loading in the data

Now import the dataset and do some processing. - The dataset has one sentence per line. - You will be doing character generation, so you have to process each sentence by converting each **character** (and not word) to a number. - You will use the ord function to convert a unique character to a unique integer ID. - Store each line in a list. - Create a data generator that takes in the batch_size and the max_length. - The max_length corresponds to the maximum length of the sentence.

```
In [ ]: dirname = 'data/'
lines = [] # storing all the lines in a variable.
for filename in os.listdir(dirname):
    with open(os.path.join(dirname, filename)) as files:
        for line in files:
            # remove leading and trailing whitespace
            pure_line = line.strip()

            # if pure_line is not the empty string,
            if pure_line:
```

```

# append it to the list
lines.append(pure_line)

```

```

In [ ]: n_lines = len(lines)
        print(f"Number of lines: {n_lines}")
        print(f"Sample line at position 0 {lines[0]}")
        print(f"Sample line at position 999 {lines[999]}")

```

Notice that the letters are both uppercase and lowercase. In order to reduce the complexity of the task, we will convert all characters to lowercase. This way, the model only needs to predict the likelihood that a letter is 'a' and not decide between uppercase 'A' and lowercase 'a'.

```

In [ ]: # go through each line
        for i, line in enumerate(lines):
            # convert to all lowercase
            lines[i] = line.lower()

        print(f"Number of lines: {n_lines}")
        print(f"Sample line at position 0 {lines[0]}")
        print(f"Sample line at position 999 {lines[999]}")

In [ ]: eval_lines = lines[-1000:] # Create a holdout validation set
        lines = lines[:-1000] # Leave the rest for training

        print(f"Number of lines for training: {len(lines)}")
        print(f"Number of lines for validation: {len(eval_lines)}")

```

1.2 Convert a line to tensor

Now that you have your list of lines, you will convert each character in that list to a number. You can use Python's ord function to do it.

Given a string representing of one Unicode character, the ord function return an integer representing the Unicode code point of that character.

```

In [ ]: # View the unique unicode integer associated with each character
        print(f"ord('a'): {ord('a')}")
        print(f"ord('b'): {ord('b')}")
        print(f"ord('c'): {ord('c')}")
        print(f"ord(' '): {ord(' ')}")
        print(f"ord('x'): {ord('x')}")
        print(f"ord('y'): {ord('y')}")
        print(f"ord('z'): {ord('z')}")
        print(f"ord('1'): {ord('1')}")
        print(f"ord('2'): {ord('2')}")
        print(f"ord('3'): {ord('3')}")

```

Exercise 01

Instructions: Write a function that takes in a single line and transforms each character into its unicode integer. This returns a list of integers, which we'll refer to as a tensor. - Use a special integer to represent the end of the sentence (the end of the line). - This will be the EOS_int (end of sentence integer) parameter of the function. - Include the EOS_int as the last integer of the - For this exercise, you will use the number 1 to represent the end of a sentence.

```

In [ ]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: line_to_tensor
        def line_to_tensor(line, EOS_int=1):
            """Turns a line of text into a tensor

            Args:
                line (str): A single line of text.
                EOS_int (int, optional): End-of-sentence integer. Defaults to 1.

            Returns:
                list: a list of integers (unicode values) for the characters in the `line`.
            """

            # Initialize the tensor as an empty list
            tensor = []
            ### START CODE HERE (Replace instances of 'None' with your code) ###
            # for each character:
            for c in line:
                # convert to unicode int
                c_int = ord(c)

                # append the unicode integer to the tensor list
                tensor.append(c_int)

            # include the end-of-sentence integer
            tensor.append(EOS_int)
            ### END CODE HERE ###

            return tensor

In [ ]: # Testing your output
        line_to_tensor('abc xyz')

```

Expected Output

```
[97, 98, 99, 32, 120, 121, 122, 1]
```

1.3 Batch generator

Most of the time in Natural Language Processing, and AI in general we use batches when training our data sets. Here, you will build a data generator that takes in a text and returns a batch of text lines (lines are sentences). - The generator converts text lines (sentences) into numpy arrays of integers padded by zeros so that all arrays have the same length, which is the length of the longest sentence in the entire data set.

Once you create the generator, you can iterate on it like this:

```
next(data_generator)
```

This generator returns the data in a format that you could directly use in your model when computing the feed-forward of your algorithm. This iterator returns a batch of lines and per token

mask. The batch is a tuple of three parts: inputs, targets, mask. The inputs and targets are identical. The second column will be used to evaluate your predictions. Mask is 1 for non-padding tokens.

Exercise 02 **Instructions:** Implement the data generator below. Here are some things you will need.

- While True loop: this will yield one batch at a time.
- if index \geq num_lines, set index to 0.
- The generator should return shuffled batches of data. To achieve this without modifying the actual lines a list containing the indexes of data_lines is created. This list can be shuffled and used to get random batches everytime the index is reset.
- if len(line) < max_length append line to cur_batch.
 - Note that a line that has length equal to max_length should not be appended to the batch.
 - This is because when converting the characters into a tensor of integers, an additional end of sentence token id will be added.
 - So if max_length is 5, and a line has 4 characters, the tensor representing those 4 characters plus the end of sentence character will be of length 5, which is the max length.
- if len(cur_batch) == batch_size, go over every line, convert it to an int and store it.

Remember that when calling np you are really calling trax.fastmath.numpy which is trax's version of numpy that is compatible with JAX. As a result of this, where you used to encounter the type numpy.ndarray now you will find the type jax.interpreters.xla.DeviceArray.

Hints

Use the line_to_tensor function above inside a list comprehension in order to pad lines with zeros.

Keep in mind that the length of the tensor is always 1 + the length of the original line of characters. Keep this in mind when setting the padding of zeros.

```
In [ ]: # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: data_generator
        def data_generator(batch_size, max_length, data_lines, line_to_tensor=line_to_tensor, shuffle=True):
            """Generator function that yields batches of data

            Args:
                batch_size (int): number of examples (in this case, sentences) per batch.
                max_length (int): maximum length of the output tensor.
                NOTE: max_length includes the end-of-sentence character that will be added
                        to the tensor.
                        Keep in mind that the length of the tensor is always 1 + the length
                        of the original line of characters.
                data_lines (list): list of the sentences to group into batches.
                line_to_tensor (function, optional): function that converts line to tensor. Defaults to line_to_tensor.
                shuffle (bool, optional): True if the generator should generate random batches. Defaults to True.

            Yields:
                tuple: two copies of the batch (jax.interpreters.xla.DeviceArray) and mask (jax.interpreters.xla.DeviceArray)"""
```

```

    NOTE: jax.interpreters.xla.DeviceArray is trax's version of numpy.ndarray
"""
# initialize the index that points to the current position in the lines index array
index = 0

# initialize the list that will contain the current batch
cur_batch = []

# count the number of lines in data_lines
num_lines = len(data_lines)

# create an array with the indexes of data_lines that can be shuffled
lines_index = [*range(num_lines)]

# shuffle line indexes if shuffle is set to True
if shuffle:
    rnd.shuffle(lines_index)

### START CODE HERE (Replace instances of 'None' with your code) ###
while True:

    # if the index is greater or equal than to the number of lines in data_lines
    if None:
        # then reset the index to 0
        index = None
        # shuffle line indexes if shuffle is set to True
        if shuffle:
            rnd.shuffle(lines_index)

    # get a line at the `lines_index[index]` position in data_lines
    line = None

    # if the length of the line is less than max_length
    if None:
        # append the line to the current batch
        cur_batch.append(None)

    # increment the index by one
    index += None

    # if the current batch is now equal to the desired batch size
    if None:

        batch = []
        mask = []

        # go through each line (li) in cur_batch
        for li in None:

```

```

    # convert the line (li) to a tensor of integers
    tensor = None

    # Create a list of zeros to represent the padding
    # so that the tensor plus padding will have length `max_length`
    pad = [0] * (None)

    # combine the tensor plus pad
    tensor_pad = None

    # append the padded tensor to the batch
    batch.append(None)

    # A mask for tensor_pad is 1 wherever tensor_pad is not
    # 0 and 0 wherever tensor_pad is 0, i.e. if tensor_pad is
    # [1, 2, 3, 0, 0, 0] then example_mask should be
    # [1, 1, 1, 0, 0, 0]
    # Hint: Use a list comprehension for this
    example_mask = None
    mask.append(example_mask)

    # convert the batch (data type list) to a trax's numpy array
    batch_np_arr = np.array(None)
    mask_np_arr = np.array(None)

    ### END CODE HERE ##

    # Yield two copies of the batch and mask.
    yield batch_np_arr, batch_np_arr, mask_np_arr

    # reset the current batch to an empty list
    cur_batch = []

```

```

In [ ]: # Try out your data generator
        tmp_lines = ['12345678901', #length 11
                     '123456789', # length 9
                     '234567890', # length 9
                     '345678901'] # length 9

        # Get a batch size of 2, max length 10
        tmp_data_gen = data_generator(batch_size=2,
                                      max_length=10,
                                      data_lines=tmp_lines,
                                      shuffle=False)

        # get one batch
        tmp_batch = next(tmp_data_gen)

```

```
# view the batch
tmp_batch
```

Expected output

```
(DeviceArray([[49, 50, 51, 52, 53, 54, 55, 56, 57, 1],
              [50, 51, 52, 53, 54, 55, 56, 57, 48, 1]], dtype=int32),
 DeviceArray([[49, 50, 51, 52, 53, 54, 55, 56, 57, 1],
              [50, 51, 52, 53, 54, 55, 56, 57, 48, 1]], dtype=int32),
 DeviceArray([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
              [1, 1, 1, 1, 1, 1, 1, 1, 1, 1]], dtype=int32))
```

Now that you have your generator, you can just call them and they will return tensors which correspond to your lines in Shakespeare. The first column and the second column are identical. Now you can go ahead and start building your neural network.

1.4 Repeating Batch generator

The way the iterator is currently defined, it will keep providing batches forever.

Although it is not needed, we want to show you the `itertools.cycle` function which is really useful when the generator eventually stops

Notice that it is expected to use this function within the training function further below

Usually we want to cycle over the dataset multiple times during training (i.e. train for multiple *epochs*).

For small datasets we can use `itertools.cycle` to achieve this easily.

```
In [ ]: import itertools
```

```
infinite_data_generator = itertools.cycle(
    data_generator(batch_size=2, max_length=10, data_lines=tmp_lines))
```

You can see that we can get more than the 5 lines in `tmp_lines` using this.

```
In [ ]: ten_lines = [next(infinite_data_generator) for _ in range(10)]
        print(len(ten_lines))
```

2 Part 2: Defining the GRU model

Now that you have the input and output tensors, you will go ahead and initialize your model. You will be implementing the GRULM, gated recurrent unit model. To implement this model, you will be using google's `trax` package. Instead of making you implement the GRU from scratch, we will give you the necessary methods from a build in package. You can use the following packages when constructing the model:

- `tl.Serial`: Combinator that applies layers serially (by function composition). [docs](#) / [source code](#)
 - You can pass in the layers as arguments to `Serial`, separated by commas.
 - For example: `tl.Serial(tl.Embeddings(...), tl.Mean(...), tl.Dense(...), tl.LogSoftmax(...))`

-
- `tl.ShiftRight`: Allows the model to go right in the feed forward. [docs](#) / [source code](#)
 - `ShiftRight(n_shifts=1, mode='train')` layer to shift the tensor to the right `n_shift` times
 - Here in the exercise you only need to specify the mode and not worry about `n_shifts`
-
- `tl.Embedding`: Initializes the embedding. In this case it is the size of the vocabulary by the dimension of the model. [docs](#) / [source code](#)
 - `tl.Embedding(vocab_size, d_feature)`.
 - `vocab_size` is the number of unique words in the given vocabulary.
 - `d_feature` is the number of elements in the word embedding (some choices for a word embedding size range from 150 to 300, for example). ____
 - `tl.GRU`: Trax GRU layer. [docs](#) / [source code](#)
 - `GRU(n_units)` Builds a traditional GRU of `n_cells` with dense internal transformations.
 - GRU paper: <https://arxiv.org/abs/1412.3555> ____
 - `tl.Dense`: A dense layer. [docs](#) / [source code](#)
 - `tl.Dense(n_units)`: The parameter `n_units` is the number of units chosen for this dense layer. ____
 - `tl.LogSoftmax`: Log of the output probabilities. [docs](#) / [source code](#)
 - Here, you don't need to set any parameters for `LogSoftMax()`. ____

Exercise 03 **Instructions:** Implement the GRULM class below. You should be using all the methods explained above.

```
In [ ]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: GRULM
        def GRULM(vocab_size=256, d_model=512, n_layers=2, mode='train'):
            """Returns a GRU language model.

            Args:
                vocab_size (int, optional): Size of the vocabulary. Defaults to 256.
                d_model (int, optional): Depth of embedding (n_units in the GRU cell). Defaults to 512.
                n_layers (int, optional): Number of GRU layers. Defaults to 2.
                mode (str, optional): 'train', 'eval' or 'predict', predict mode is for fast inference.

            Returns:
                trax.layers.combinators.Serial: A GRU language model as a layer that maps from input to output.
            """

            ### START CODE HERE (Replace instances of 'None' with your code) ###
            model = tl.Serial(
                None, # Stack the ShiftRight layer
```

```

        None, # Stack the embedding layer
        None, # Stack GRU layers of d_model units keeping n_layer parameter in mind (use
        None, # Dense layer
        None # Log Softmax
    )
    ### END CODE HERE ###
    return model

In [ ]: # testing your model
        model = GRULM()
        print(model)

```

Expected output

```

Serial[
  ShiftRight(1)
  Embedding_256_512
  GRU_512
  GRU_512
  Dense_256
  LogSoftmax
]

```

Part 3: Training

Now you are going to train your model. As usual, you have to define the cost function, the optimizer, and decide whether you will be training it on a gpu or cpu. You also have to feed in a built model. Before, going into the training, we re-introduce the `TrainTask` and `EvalTask` abstractions from the last week's assignment.

To train a model on a task, Trax defines an abstraction `trax.supervised.training.TrainTask` which packages the train data, loss and optimizer (among other things) together into an object.

Similarly to evaluate a model, Trax defines an abstraction `trax.supervised.training.EvalTask` which packages the eval data and metrics (among other things) into another object.

The final piece tying things together is the `trax.supervised.training.Loop` abstraction that is a very simple and flexible way to put everything together and train the model, all the while evaluating it and saving checkpoints. Using `training.Loop` will save you a lot of code compared to always writing the training loop by hand, like you did in courses 1 and 2. More importantly, you are less likely to have a bug in that code that would ruin your training.

```

In [ ]: batch_size = 32
        max_length = 64

```

An epoch is traditionally defined as one pass through the dataset.

Since the dataset was divided in batches you need several steps (gradient evaluations) in order to complete an epoch. So, one epoch corresponds to the number of examples in a batch times the number of steps. In short, in each epoch you go over all the dataset.

The `max_length` variable defines the maximum length of lines to be used in training our data, lines longer than that length are discarded.

Below is a function and results that indicate how many lines conform to our criteria of maximum length of a sentence in the entire dataset and how many steps are required in order to cover the entire dataset which in turn corresponds to an epoch.

```
In [ ]: def n_used_lines(lines, max_length):
        '''
        Args:
        lines: all lines of text an array of lines
        max_length - max_length of a line in order to be considered an int
        Return:
        n_lines -number of efective examples
        '''

        n_lines = 0
        for l in lines:
            if len(l) <= max_length:
                n_lines += 1
        return n_lines

num_used_lines = n_used_lines(lines, 32)
print('Number of used lines from the dataset:', num_used_lines)
print('Batch size (a power of 2):', int(batch_size))
steps_per_epoch = int(num_used_lines/batch_size)
print('Number of steps to cover one epoch:', steps_per_epoch)
```

Expected output:

Number of used lines from the dataset: 25881

Batch size (a power of 2): 32

Number of steps to cover one epoch: 808

3.1 Training the model

You will now write a function that takes in your model and trains it. To train your model you have to decide how many times you want to iterate over the entire data set.

Exercise 04

Instructions: Implement the `train_model` program below to train the neural network above. Here is a list of things you should do:

- Create a `trax.supervised.trainer.TrainTask` object, this encapsulates the aspects of the dataset and the problem at hand:
 - `labeled_data` = the labeled data that we want to *train* on.
 - `loss_fn` = `tl.CrossEntropyLoss()`
 - `optimizer` = `trax.optimizers.Adam()` with learning rate = 0.0005
- Create a `trax.supervised.trainer.EvalTask` object, this encapsulates aspects of evaluating the model:
 - `labeled_data` = the labeled data that we want to *evaluate* on.
 - `metrics` = `tl.CrossEntropyLoss()` and `tl.Accuracy()`
 - How frequently we want to evaluate and checkpoint the model.

- Create a `trax.supervised.trainer.Loop` object, this encapsulates the following:
 - The previously created `TrainTask` and `EvalTask` objects.
 - the training model = Section ??
 - optionally the evaluation model, if different from the training model. NOTE: in presence of Dropout etc we usually want the evaluation model to behave slightly differently than the training model.

You will be using a cross entropy loss, with Adam optimizer. Please read the [trax](#) documentation to get a full understanding. Make sure you use the number of steps provided as a parameter to train for the desired number of steps.

NOTE: Don't forget to wrap the data generator in `itertools.cycle` to iterate on it for multiple epochs.

```
In [ ]: from trax.supervised import training

# UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: train_model
def train_model(model, data_generator, batch_size=32, max_length=64, lines=lines, eval_lines=eval_lines, n_steps=1, output_dir=None):
    """Function that trains the model

    Args:
        model (trax.layers.combinators.Serial): GRU model.
        data_generator (function): Data generator function.
        batch_size (int, optional): Number of lines per batch. Defaults to 32.
        max_length (int, optional): Maximum length allowed for a line to be processed.
        lines (list, optional): List of lines to use for training. Defaults to lines.
        eval_lines (list, optional): List of lines to use for evaluation. Defaults to eval_lines.
        n_steps (int, optional): Number of steps to train. Defaults to 1.
        output_dir (str, optional): Relative path of directory to save model. Defaults to None.

    Returns:
        trax.supervised.training.Loop: Training loop for the model.
    """

    ### START CODE HERE (Replace instances of 'None' with your code) ###
    bare_train_generator = None
    infinite_train_generator = None

    bare_eval_generator = None
    infinite_eval_generator = None

    train_task = training.TrainTask(
        labeled_data=None, # Use infinite train data generator
        loss_layer=None, # Don't forget to instantiate this object
        optimizer=None # Don't forget to add the learning rate parameter
    )

    eval_task = training.EvalTask(
```

```

        labeled_data=None,      # Use infinite eval data generator
        metrics=[None, None],  # Don't forget to instantiate these objects
        n_eval_batches=3       # For better evaluation accuracy in reasonable time
    )

    training_loop = training.Loop(model,
                                  train_task,
                                  eval_task=eval_task,
                                  output_dir=output_dir)

    training_loop.run(n_steps=None)

    ### END CODE HERE ###

    # We return this because it contains a handle to the model, which has the weights
    return training_loop

```

```

In [ ]: # Train the model 1 step and keep the `trax.supervised.training.Loop` object.
        training_loop = train_model(GRULM(), data_generator)

```

The model was only trained for 1 step due to the constraints of this environment. Even on a GPU accelerated environment it will take many hours for it to achieve a good level of accuracy. For the rest of the assignment you will be using a pretrained model but now you should understand how the training can be done using Trax.

Part 4: Evaluation

4.1 Evaluating using the deep nets

Now that you have learned how to train a model, you will learn how to evaluate it. To evaluate language models, we usually use perplexity which is a measure of how well a probability model predicts a sample. Note that perplexity is defined as:

$$P(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i|w_1, \dots, w_{n-1})}}$$

As an implementation hack, you would usually take the log of that formula (to enable us to use the log probabilities we get as output of our RNN, convert exponents to products, and products into sums which makes computations less complicated and computationally more efficient). You should also take care of the padding, since you do not want to include the padding when calculating the perplexity (because we do not want to have a perplexity measure artificially good).

$$\begin{aligned}
 \log P(W) &= \log \left(\sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i|w_1, \dots, w_{n-1})}} \right) \\
 &= \log \left(\prod_{i=1}^N \frac{1}{P(w_i|w_1, \dots, w_{n-1})} \right)^{\frac{1}{N}} \\
 &= \log \left(\prod_{i=1}^N P(w_i|w_1, \dots, w_{n-1}) \right)^{-\frac{1}{N}}
 \end{aligned}$$

$$= -\frac{1}{N} \log \left(\prod_{i=1}^N P(w_i | w_1, \dots, w_{n-1}) \right)$$

$$= -\frac{1}{N} \left(\sum_{i=1}^N \log P(w_i | w_1, \dots, w_{n-1}) \right)$$

Exercise 05 **Instructions:** Write a program that will help evaluate your model. Implementation hack: your program takes in preds and target. Preds is a tensor of log probabilities. You can use `tl.one_hot` to transform the target into the same dimension. You then multiply them and sum.

You also have to create a mask to only get the non-padded probabilities. Good luck!

Hints

To convert the target into the same dimension as the predictions tensor use `tl.one_hot` with target and `preds.shape[-1]`.

You will also need the `np.equal` function in order to unpad the data and properly compute perplexity.

Keep in mind while implementing the formula above that w_i represents a letter from our 256 letter alphabet.

```
In [ ]: # UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # GRADED FUNCTION: test_model
        def test_model(preds, target):
            """Function to test the model.

            Args:
                preds (jax.interpreters.xla.DeviceArray): Predictions of a list of batches of
                target (jax.interpreters.xla.DeviceArray): Actual list of batches of tensors c

            Returns:
                float: log_perplexity of the model.
            """
            ### START CODE HERE (Replace instances of 'None' with your code) ###
            total_log_ppx = np.sum(None * None, axis=-1) # HINT: tl.one_hot() should replace

            non_pad = 1.0 - np.equal(None, None)         # You should check if the target equ
            ppx = None * None                             # Get rid of the padding

            log_ppx = np.sum(None) / np.sum(None)
            ### END CODE HERE ###

            return -log_ppx
```

```
In [ ]: # UNQ_C6 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        # Testing
        model = GRULM()
        model.init_from_file('model.pkl.gz')
        batch = next(data_generator(batch_size, max_length, lines, shuffle=False))
        preds = model(batch[0])
```

```
log_ppx = test_model(preds, batch[1])
print('The log perplexity and perplexity of your model are respectively', log_ppx, np.
```

Expected Output: The log perplexity and perplexity of your model are respectively around 1.9 and 7.2.

Part 5: Generating the language with your own model

We will now use your own language model to generate new sentences for that we need to make draws from a Gumble distribution.

The Gumbel Probability Density Function (PDF) is defined as:

$$f(z) = \frac{1}{\beta} e^{(-z + e^{(-z)})}$$

where:

$$z = \frac{(x - \mu)}{\beta}$$

The maximum value, which is what we choose as the prediction in the last step of a Recursive Neural Network RNN we are using for text generation, in a sample of a random variable following an exponential distribution approaches the Gumbel distribution when the sample increases asymptotically. For that reason, the Gumbel distribution is used to sample from a categorical distribution.

```
In [ ]: # Run this cell to generate some news sentence
def gumbel_sample(log_probs, temperature=1.0):
    """Gumbel sampling from a categorical distribution."""
    u = numpy.random.uniform(low=1e-6, high=1.0 - 1e-6, size=log_probs.shape)
    g = -np.log(-np.log(u))
    return np.argmax(log_probs + g * temperature, axis=-1)

def predict(num_chars, prefix):
    inp = [ord(c) for c in prefix]
    result = [c for c in prefix]
    max_len = len(prefix) + num_chars
    for _ in range(num_chars):
        cur_inp = np.array(inp + [0] * (max_len - len(inp)))
        outp = model(cur_inp[None, :]) # Add batch dim.
        next_char = gumbel_sample(outp[0, len(inp)])
        inp += [int(next_char)]

    if inp[-1] == 1:
        break # EOS
    result.append(chr(int(next_char)))

    return "".join(result)

print(predict(32, ""))

In [ ]: print(predict(32, ""))
print(predict(32, ""))
print(predict(32, ""))
```

In the generated text above, you can see that the model generates text that makes sense capturing dependencies between words and without any input. A simple n-gram model would have not been able to capture all of that in one sentence.

On statistical methods

Using a statistical method like the one you implemented in course 2 will not give you results that are as good. Your model will not be able to encode information seen previously in the data set and as a result, the perplexity will increase. Remember from course 2 that the higher the perplexity, the worse your model is. Furthermore, statistical ngram models take up too much space and memory. As a result, it will be inefficient and too slow. Conversely, with deepnets, you can get a better perplexity. Note, learning about n-gram language models is still important and allows you to better understand deepnets.